


CS3485

Deep Learning for Computer Vision



Lec 16: Advanced GANs

Announcements

- Teresa was born! Thank you for your patience on Monday!



- I made a few changes for our next lectures (see [schedule](#) on our website).


Lecture	11/08/2023 Wednesday	Lec17 - Advanced GANs
Lecture	11/13/2023 Monday	Lec18 - Transformers and ChatGPT
Lecture	11/15/2023 Wednesday	Lec19 - Image Generation by Prompt
Review Session	11/17/2023 Friday	Second Review Session
Exam	11/20/2023 Monday	Final Exam
	11/22/2023 Wednesday	THANKSGIVING BREAK!

Announcements


- Project Proposal:
 - Due on Nov 10th, and there is a submission link on canvas,
 - Remember it counts as part of the grade!
- Info about late submissions on the website (more for next year, actually).
- Lab 6:
 - Students were confused, so I added a hint to the instructions about the MSE loss.
 - Due tonight (but I am willing to give one extra day if you ask, because of the above).
- Lab 7:
 - Will be an extra lab: it will be worth 30 points to be redistributed (added) on the grades of the previous labs, starting from the lowest grade.
 - Does not need to be in teams. Can be individual.
 - It will be due the last day of classes (Dec 8th). No late days points are allowed, even if you have some leftover.
- Interesting application of dense pose estimation.

(Tentative) Lecture Roadmap

Basics of Deep Learning



Computer Vision Tasks



More interesting GANs


- Last time we saw how GANs can generate new digits from the MNIST dataset and new faces.
- Although interesting, these results **were not realist enough** compared to more modern GAN architectures.
- Today, we'll see how modern GANs (such as StyleGAN) are able to generate **visually stunning high-resolution face images!**
- Before that, we'll also see how to **conditionally generate new images** using GANs which will provide us with tools to solve many other problems in image generation.



New faces generated by StyleGAN

Conditional GANs

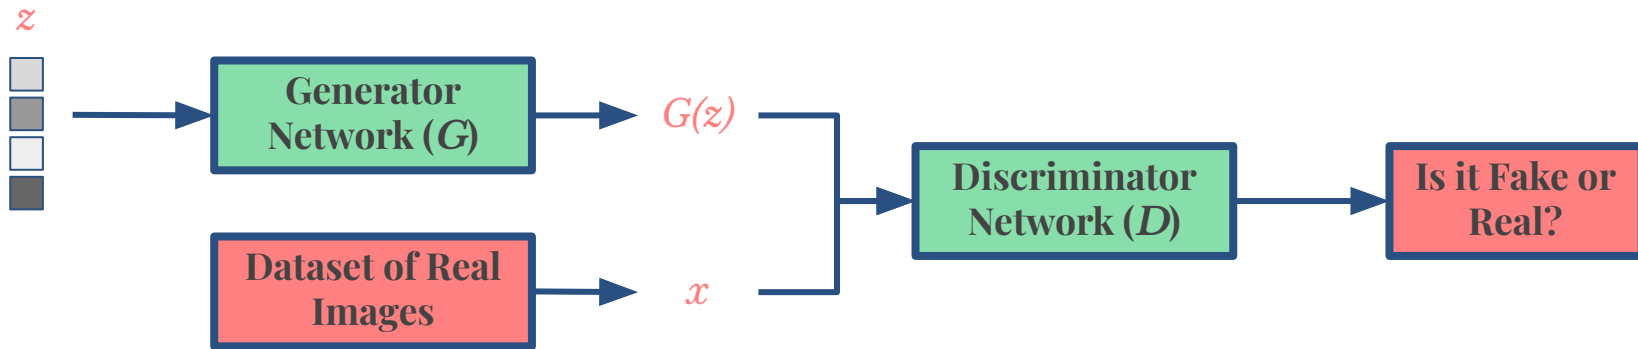
- All GAN models we have seen so far model a probability density in high dimension and provide means to sample according to it, which is useful **for image synthesis only**.
- However, most of the practical applications require the ability to sample a **conditional distribution**, i.e., sample new data conditioned on some information we have at our disposal.
- For example, we may want to sample a datapoint conditioned on its class (I may want to sample only new 7's instead of any random digit).
- **Conditional GAN**, [published](#) in 2014, was conceived to adapt our previous, simple GAN architecture (called **Vanilla GAN**) to this setting.



New MNIST digits generated according to their classes.

Conditional GANs

- Let's first review our previous GAN approach:
 - We have a **Generator Network** G that takes in a random vector z and produces a new, generated image $G(z)$.
 - We also have a **Discriminator Network** D that takes in an image as its input and classifies it in fake (i.e., generated by G) or real (i.e., coming from an image dataset).
 - The goal is twofold: (1) train a very good discriminator network and (2) train a generator that beats this discriminator.



Conditional GANs

- In Conditional GAN, the same training approach is taken, but now both generator and discriminator inputs will carry **class information**.
- To do that, we just need to “add more data” to both inputs. Say we have K classes ($K = 10$ for MNIST):
 - For the **Generator input**, append to z a vector of K dimensional one-hot encoding of the class you want the generated image to be from.
 - For the **Discriminator input**, append K more channels to the input image such that they work as an one-encoding of the that image’s class (from either the image dataset or the generator’s input)*.

* Note that, even if the image is realistic, but the class that image is attached to is not the correct one the discriminator here should output “fake”.

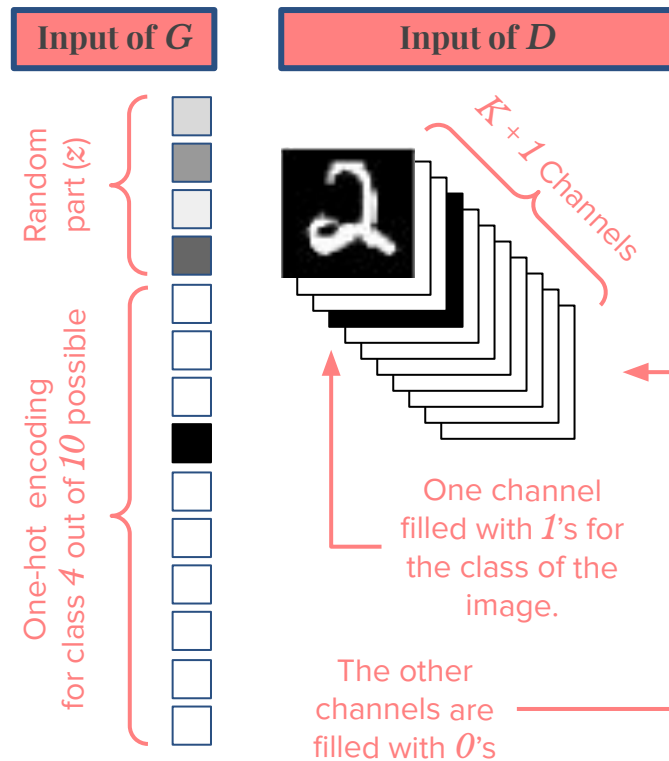


Image-to-Image Translation

- We can use the same principle of conditioning the generation to a class to create interesting conditions.
- For example, we may want to generate a realistic image conditioned in a certain edge map, i.e., a new image that has its edges given by the user.
- This approach will be very useful for the task of **Image-to-Image Translation**:

Image-to-image translation is the task of taking images from one domain and transforming them so they have the characteristics of images from another domain.

- In our example above, we converted an image in the domain of edges to the domain of realistic RGB images.

Edge Map



Edge Source



New Image




Pix2Pix

- Published in 2016, the Pix2Pix strategy to solve image to image translation involved a GAN network that used the concepts from Conditional GANs.
- Here, the difference is that the generator receives an image input z in one domain (edge map, for example) and outputs the corresponding image on the other domain.
- The discriminator is then tasked to check if the pairings edge map/image are realistic.

Training on fake data



Training on real data



Pix2Pix

Edges to Photos

Input




Ground Truth



Output



Input



Ground Truth



Output



Pix2Pix

- Note that the edge maps don't need to be realistic. These are the results from when you input a line drawing to generator trained on edge maps.

Drawings to Photo



Pix2Pix Applications


- Pix2Pix has been applied in translation domains beyond that of edges to RGB images (but always following the same training strategy).
- Here, you can have the generator generate aerial photos from a map or maps from aerial photos.



Pix2Pix Applications

- In a similar way, Pix2Pix was used to generate new building facades according to a image of facade labels, i.e., positions of windows, doors, roofs, etc.

Facade labels to Photo

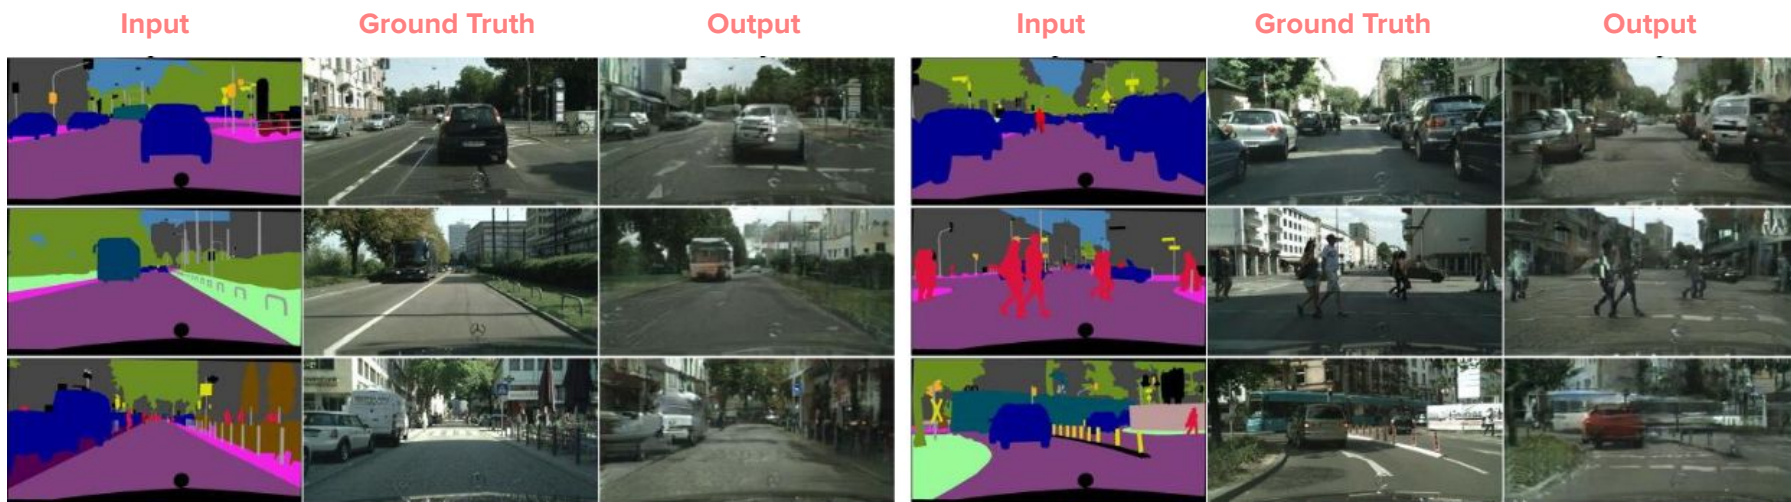


- You can actually try out some of these algorithms yourself! In this [link](#), you'll find the edge to image and the facade labels to image applications.

Pix2Pix Applications

- Pix2Pix can be applied to image generation conditioned on a given semantic segmentation.

Segmentation to Photo




Pix2Pix Applications

- The principles of Pix2Pix have also been applied to many artistic endeavors.

Learning to see work piece



GauGAN art generator



Pix2Pix Applications

Day Image to Night Image

Input



Ground Truth



Output




Input



Ground Truth




Output



Thermal Image to Photo

Input



Ground Truth



Output



Input



Ground Truth



Output



Pix2Pix Applications

Season changer



winter Yosemite → summer Yosemite



summer Yosemite → winter Yosemite



Pix2Pix Applications

Photo Enhancement (post-hoc focusing) and Painting Style Transfer

Input



Output



Input



Output



Input




Output




Input




Monet




Van Gogh



Cezanne



Ukyo-e



Exercise (*in pairs*)

- Play with Pix2Pix! You can go to this [link](#) and try out some of their algorithms. What do you notice when you play with it?

Getting high resolution images

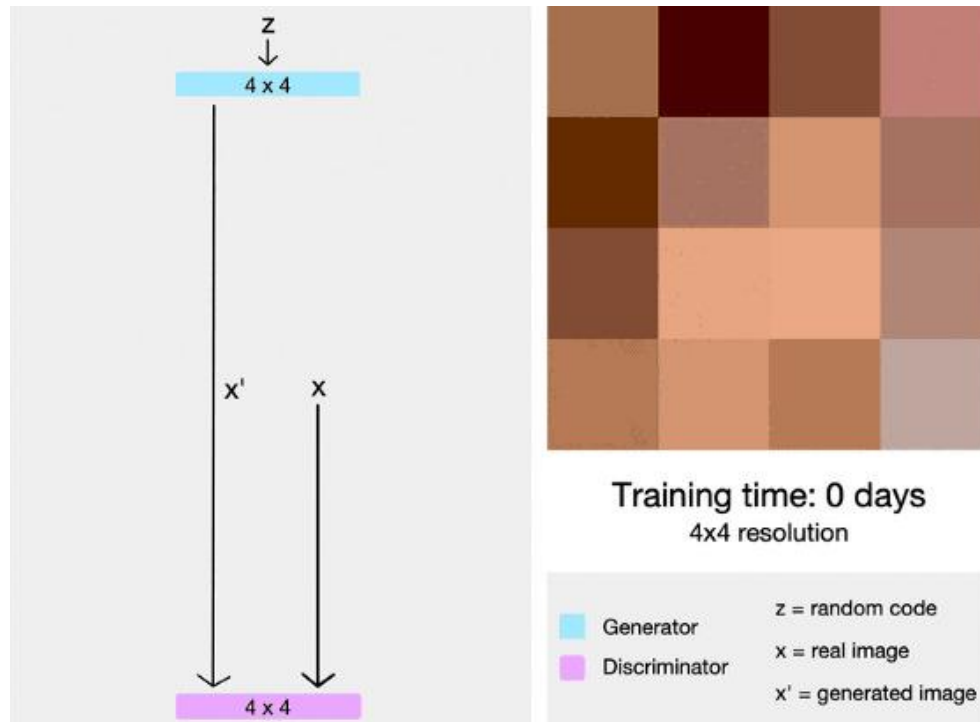
- We saw that using GANs we can generate small images in various settings, but how can one generate **high resolution images** (images that contain fine level visual features)?
- Standard GANs could work here, but they would not be practical to generate high quality images (1024×1024 size) because of their architecture limitations.
- The first attempt to solve this issue was proposed in 2017 and was called **ProGAN** (**Progressive Generative Adversarial Networks**).



Generated face using ProGAN.

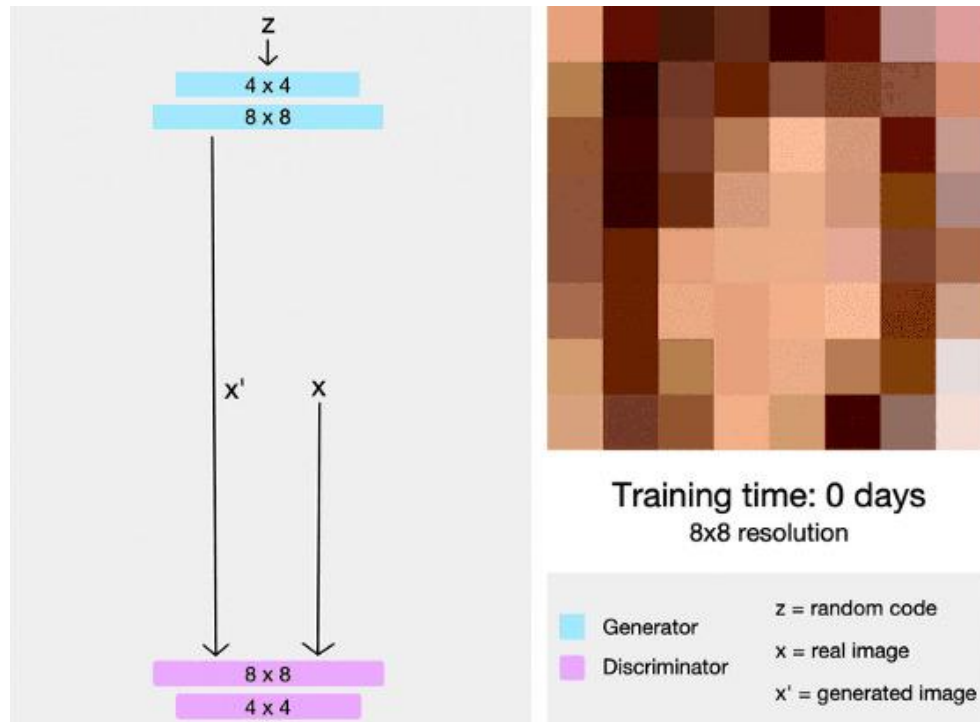
ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.



ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.



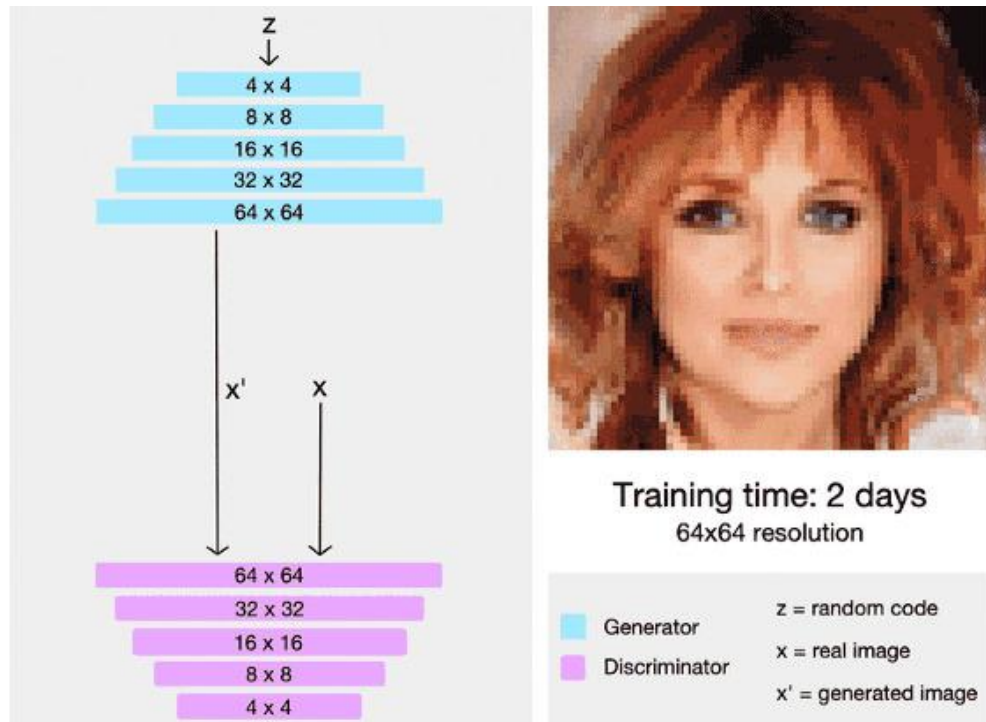
ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.



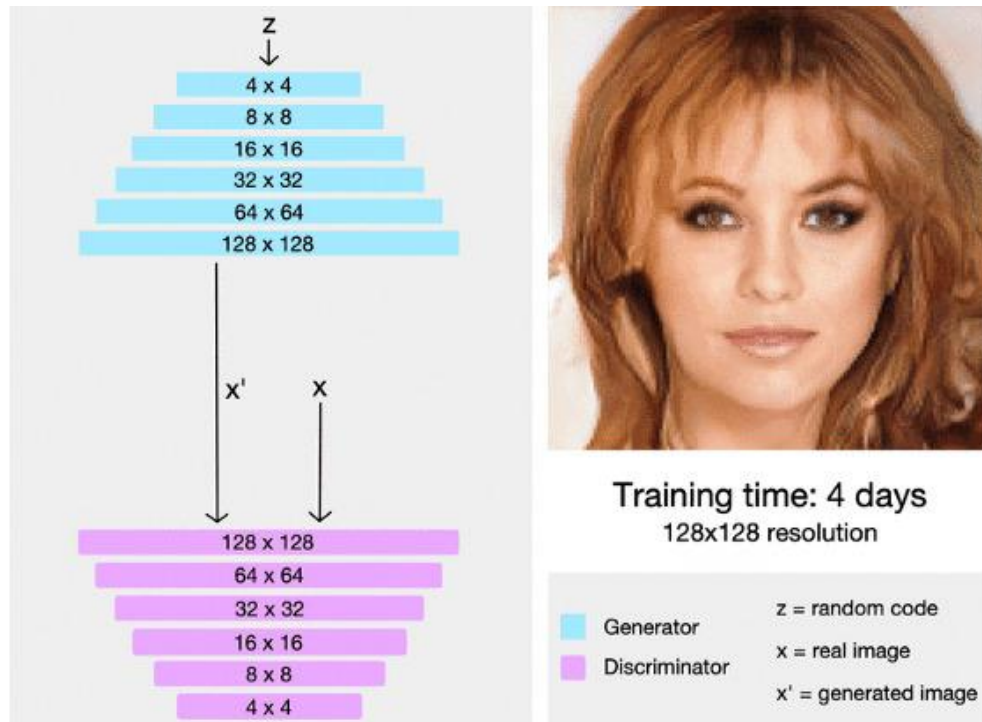
ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.




Training time: 6 days
256x256 resolution




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.





Training time: 7 days
256x256 resolution




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.





Training time: 10 days
512x512 resolution




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.




Training time: 12 days
1024x1024 resolution




ProGAN

- ProGAN is based on an efficient way (in terms of training time) to train a GAN for High Res images.
- Instead of attempting to train all layers of the generator and discriminator at once, ProGAN trains them one layer at a time, to learn progressively higher resolution versions of the images.
- When the images generated in given resolution are good enough, we proceed to the next resolution.




Training time: 14 days
1024x1024 resolution




ProGAN

- The progressive growth in ProGAN time allowed the training on much bigger datasets of very large images in a much quicker time compared to when the layers were fixed.
- Although ProGAN expanded vanilla GANs ability to generate high-res images, still lacked the control over the **styling** of the output.
- This means that we couldn't change specific features such pose, face shape and hairstyle in a generated image from ProGAN.
- Considering this issue, the same ProGAN authors proposed **StyleGAN** in 2018.



StyleGAN

- StyleGAN mainly improves upon the existing architecture of Generator network to achieve the desired results and keeps Discriminator network and everything else **untouched**.

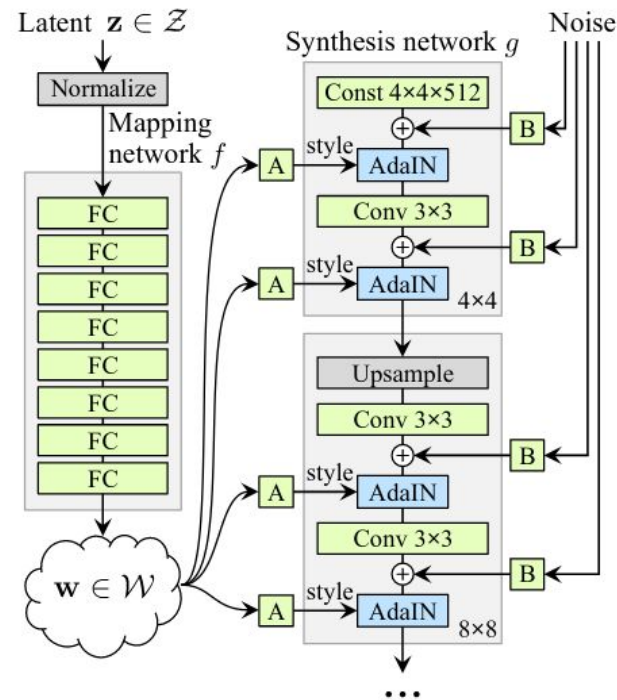


These “PixelNorm” layer is similar to Batch Normalization: It normalizes the feature vector in each pixel to unit length, and is applied after the convolutional layers in the generator. This is done to prevent signal magnitudes from spiraling out of control during training.

Generator in ProGAN

StyleGAN

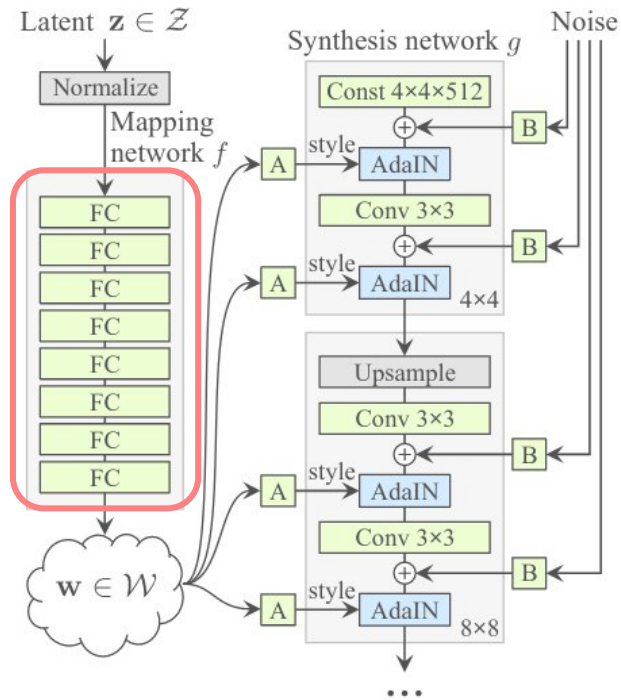
- StyleGAN mainly improves upon the existing architecture of Generator network to achieve the desired results and keeps Discriminator network and everything else **untouched**.



Generator in StyleGAN.

StyleGAN

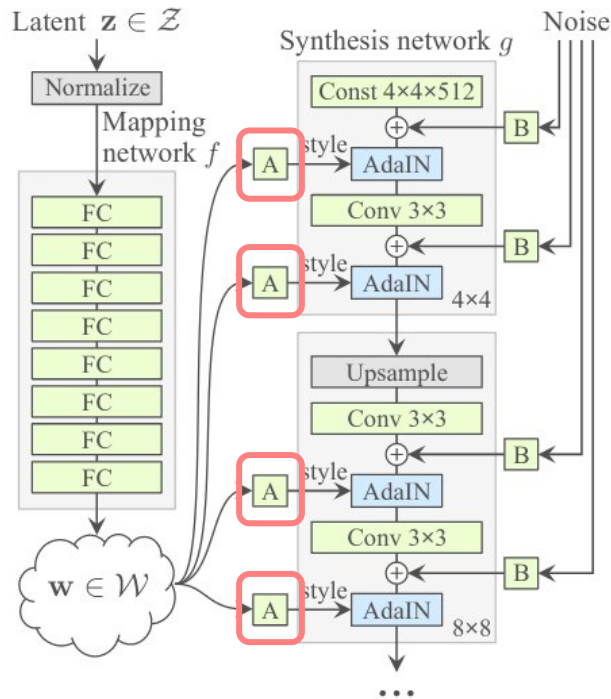
- StyleGAN mainly improves upon the existing architecture of Generator network to achieve the desired results and keeps Discriminator network and everything else **untouched**.
- The new generator has the following novelties:
 - The latent vector \mathbf{z} is first transformed into what is called a vector $\mathbf{w} = f(\mathbf{z})$ via a mapping network f .



Generator in StyleGAN.

StyleGAN

- StyleGAN mainly improves upon the existing architecture of Generator network to achieve the desired results and keeps Discriminator network and everything else **untouched**.
- The new generator has the following novelties:
 - The latent vector \mathbf{z} is first transformed into what is called a vector $\mathbf{w} = f(\mathbf{z})$ via a mapping network f .
 - \mathbf{w} is sent to MLPs (named “A” on the right, two per resolution level) that output the **style** $\mathbf{y} = (y_s, y_b)$.

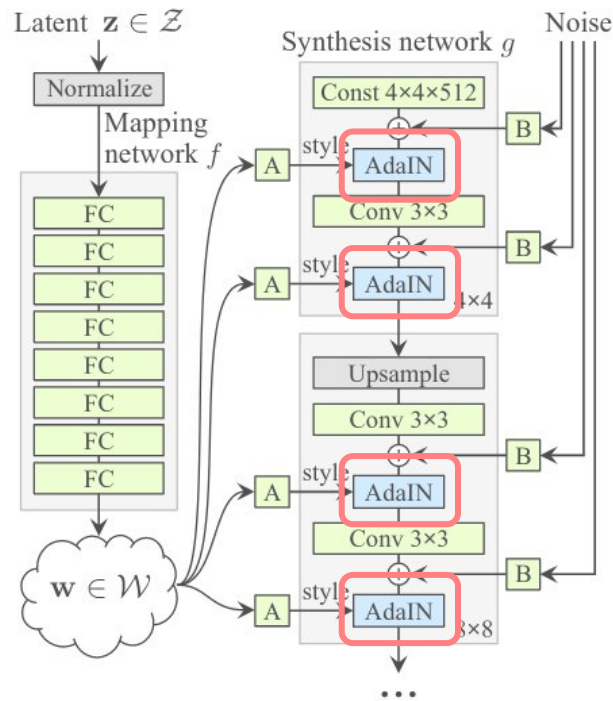


Generator in StyleGAN.

StyleGAN

- StyleGAN mainly improves upon the existing architecture of Generator network to achieve the desired results and keeps Discriminator network and everything else **untouched**.
- The new generator has the following novelties:
 - The latent vector \mathbf{z} is first transformed into what is called a vector $\mathbf{w} = f(\mathbf{z})$ via a mapping network f .
 - \mathbf{w} is sent to MLPs (named “A” on the right, two per resolution level) that output the **style** $\mathbf{y} = (y_s, y_b)$.
 - On the **synthesis network** g , a learned constant tensor gets sequentially mixed with each level’s style \mathbf{y} via an AdaIN operation in order to generate a full size image.

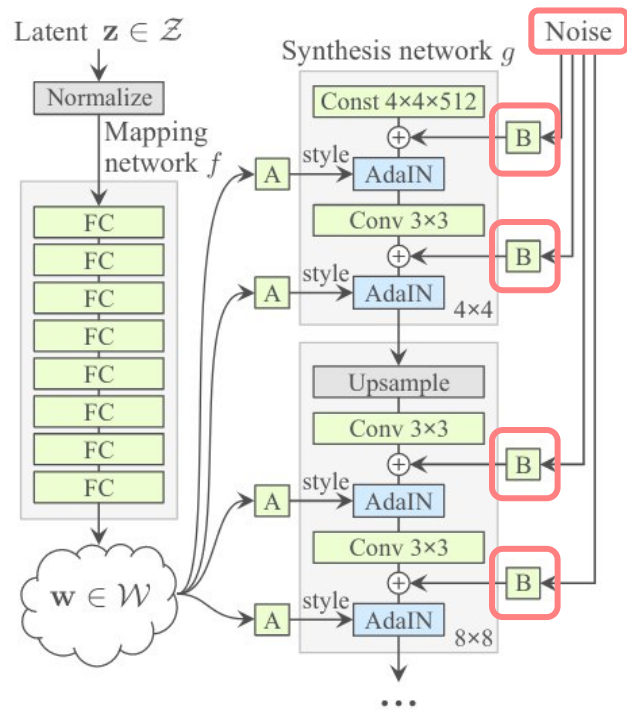
$$\text{AdaIN}(\mathbf{x}_i, \mathbf{y}) = \mathbf{y}_{s,i} \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{b,i},$$



Generator in StyleGAN.

StyleGAN

- StyleGAN mainly improves upon the existing architecture of Generator network to achieve the desired results and keeps Discriminator network and everything else **untouched**.
- The new generator has the following novelties:
 - The latent vector z is first transformed into what is called a vector $w = f(z)$ via a mapping network f .
 - w is sent to MLPs (named “A” on the right, two per resolution level) that output the **style** $y = (y_s, y_b)$.
 - On the **synthesis network** g , a learned constant tensor gets sequentially mixed with each level’s style y via an AdaIN operation in order to generate a full size image.
 - Finally, **noise** is inserted into g via some MLPs (the “B”s) to introduce style variation at a given level of detail.



Generator in StyleGAN.


Mixing styles in StyleGAN

- With StyleGAN we can mix the styles of different generated images!
- Here, two sets of images were generated from their respective latent codes (sources A and B).
- The other images were generated by copying a subset of styles y from B and the rest from A.
- Coarse y 's are those from 4^2 and 8^2 resolutions.



Mixing styles in StyleGAN

- Middle and fine y 's are from resolutions $16^2 - 32^2$ and $64^2 - 1024^2$, resp.
- Here we note that:
 - Coarse y 's correspond to high-level aspects such as general hair style pose, face shape...
 - Middle y 's relate to small facial aspects, hair style, eyes open/closed.
 - Middle y 's brings mainly the color scheme and microstructure.



Training of StyleGAN and other versions

- In the StyleGAN paper, the authors also introduce a new dataset of human faces called Flickr-Faces-HQ Dataset (FFHQ) consisting of 70,000 high-quality face images with which they trained their networks.



- StyleGAN was improved in a few ways in StyleGAN2 ([published](#) in 2019) and StyleGAN3 ([published](#) in 2021). Their main contributions are related to removing weird unexpected generated artifacts and make styles be learned in a more natural hierarchical manner.
- A nice thing about StyleGANs: their codes are available online ([here](#), [here](#) and [here](#)) and many people trained them in other datasets and released the models ([here](#) and [here](#))!

Applications of StyleGAN

- The ability to generate some many high fidelity controllable face generation has sparked many applications (for the good and for the bad). Some of them are:

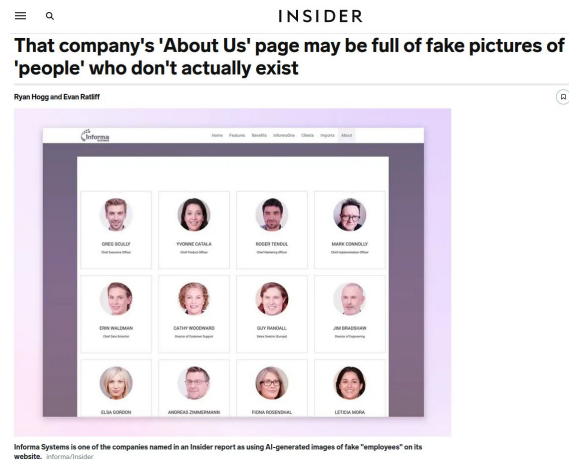
Face Interpolation



Image Editing



Well, Generate Faces (duh!)



Exercise (*in pairs*)

- The same concept of StyleGAN has been applied to many of image domains other than faces (cats, horses, memes...). [Here](#) is a website of a collection of artificially generated images from various domains (unfortunately, some of the links are broken, [here](#) is a link for face generation). Play around with them!

Video: *Bringing dead people back to life*

